

# Algorithm for Process p<sub>1</sub>

```
while (true) {  
    flag[i] = true  
    turn = j  
    while (flag[j] && turn == j) {  
        // critical section  
        flag[i] = false  
        //remainder section  
    }  
}
```

## Peterson's Solution

## Synchronization Hardware

- Many sys provide HW support for implementing the critical section code.
- Uniprocessors
  - Code would execute wi/out preemption
  - Too inefficient on multi-processor sys
    - OSs using this not scalable
- Hardware support
  - Memory barriers
  - Hardware instructions
  - Atomic variables

Hardware	SW + HW
Disable interrupts	mutex lock
Atomic instructions	semaphonre
-- test_and_set	monitor
-- compare_and_swap	

Users cannot block interrupts - must be kernel

No I/O inside critical section code - don't block anything. Can have a ContextSwitch though

Atomic: There may be a ContextSwitch, there may not be.

## Memory Barriers

- An instruction that forces any change in memory to be propagated to all other processes
- Memory model: mem guarantees a comp. arch. makes to application programs
  - Strongly ordered: all mem mods of one processors is immediately visible to all other processes
  - Weakly ordered: all mem mods of one process may not be immediately visible to all other processes
- Could add a mem barrier to the following to ensure thread 1 outputs 100:  
t1 now does:

```
while(!flag):  
    mem_barrier()  
print x
```

t2 now does:

```
x = 100  
mem_barrier()  
flag = true
```

## Hardware Instructions

Special HW instructions that *test-and-modify* the content of a word, or *swap* the contents of two words

- **Test-and-set** instruction
- **Compare and swap** instruction

### test\_and\_set instruction

```
bool test_and_set(boolean* target) {  
    bool rv = *target  
    *target = true  
    return rv  
}
```

- Executed atomically
- Returns the og value of the passed param
- Set the new val of passed param to true

## compare\_and\_swap

```
int compare_and_swap(int* value, int expected, int new_val) {
    int temp = *value

    if (*value == expected)
        *value = new_value
    return temp
}
```

- exec atomicaaly
- Return the og val of the passes param val
- Set the var val the val of the passed param new\_val

## Solution using test\_and\_swap

- Shared boolean variable lock , intialized to false
- Solution:

```
do {
    while test_and_set(&lock)
        // do nothing
        // crictical section
        lock = false
        // remainder section
    } while true
```

Critical section problem:

- Satisfies Mutual exclusion problem
- Satisfied Progress problem
- Does NOT satisfy Bounded waiting problem
  - One process can grab the same lock over and over again
  - Process can starve.
  - Will never enter critical section because it can never lock
  - BUT, there's no deadlock or livelock

## Solution using compare\_and\_swap

Shared int lock initialized to 0

```
while true:
    while compare_and_swap(&lock, 0, 1) != 0:
        //do nothing

    // critical section
    lock = 0
    //remainder section
```

## Bounded waiting mutual-exclusion with compare\_and\_swap

```
while true:
    waiting[i] = true
    key = 1
    while waiting[i] and key == 1:
        key = compare_and_swap(&lock, 0, 1)
    waiting[i] = false

    //critical section

    j = (i+1) % n
    while (j != i) and !waiting[j]:
        j = (j + 1) % n

    if j == i:
        lock = 0
    else:
        waiting[j] = false

    //remainder section
```

- Solves critical section problem
- Satisfies bounded waiting problem
- Has to wait (at most) for n-1 processes to grab the lock. Checks processes in a round-robin way to grant lock.

## Atomic Variables

- Typically, instructions such as `compare_and_swap` are used as building blocks for other synchronization tools

- One tool is an atomic variable that provides atomic (uninterruptible) updates on basic data types such as ints and bools
- ex: the `increment()` operation on the atomic var `sequence` ensures `sequence` is incremented without interruption  
`increment(&sequence)`

## Consumer

```
while true:
    while counter == 0
        // do nothing
    next_consumed = buffer[out]
    out = (out + 1) % BUFFER_SIZE
    counter--
    //consume the item in next consumed
```

## Race Conditions

- `counter++`  
could be: `reg1 = count; reg1 = reg1 + 1; count = reg1`
- `counter--`  
could be: `reg2 = count; reg2 = reg2 - 1; count = reg2`
- Consider:  
s0 producer exec `reg1 = counter`  
s1 producer exec `reg1`
- The `increment()` function can be implemented as:

```
void increment(atomic_int* v):
    int temp
    do:
        temp = *v
    while temp != compare_and_swap(v, temp, temp+1)
```

## Mutex Locks

- Previous solutions are complicated and generally inaccessible to application programmers
- OS designers build software tools to solve critical section problem
- Simplest is Mutex Lock
- Protect a C.S by first `acquire()` a lock, then `release()` the lock

- bool var indicating if lock is available or not
- Call to `acquire()` and `release()` must be atomic
  - Usually implemented via hardware atomic instructions such as `compare_and_swap`
- But this solution requires **Busy waiting**
  - This lock is therefore called a **spinlock**

## Definitions

```
acquire()
while !available
    // busy wait
available = false
```

```
release()
available = true
```

these two funcs must be implemented atomically. Both `test_and_set` and `compare_and_swap` can be used to implement these functions.

## Semaphore

- Synchronization tool that provides more sophisticated ways (than mutex locks) for process to sync their activities
- Semaphore `S` - integer var
- Can only be accessed via two indivisible (atomic) operations
  - `wait()` and `signal()`
    - (originally called `P()` and `V()`)
- Definition of `wait()` and `signal()` operations

```
wait(S)
while S <= 0
    // busy wait
S--
```

```
signal(S):
S++
```

# Midterm - moved from Oct 22nd to Oct 29th